

# GuardrailProbe Independent Benchmark — June 2026

|                 |                                                                                                  |
|-----------------|--------------------------------------------------------------------------------------------------|
| Run ID          | aaf34a3a-b850-4085-b01a-4fbcd70983ee                                                             |
| Timestamp       | 2026-06-24T12:30:56.708934+00:00                                                                 |
| Generated       | 2026-06-24T12:32:56Z                                                                             |
| Backends tested | nemo, guardrails_ai, presidio, lakera, openai_moderation, azure_content_safety, azure_prompt_shi |
| Best overall    | nemo                                                                                             |
| Worst overall   | azure_content_safety                                                                             |

## Backend Comparison Summary

| Backend              | Total | Passed | Failed | Pass % | Avg ms |
|----------------------|-------|--------|--------|--------|--------|
| nemo ★               | 78    | 67     | 11     | 85.9%  | 916    |
| guardrails ai        | 78    | 2      | 76     | 2.6%   | 87     |
| presidio             | 78    | 5      | 73     | 6.4%   | 1025   |
| lakera               | 78    | 65     | 13     | 83.3%  | 458    |
| openai moderation    | 1     | 1      | 0      | 100.0% | 34453  |
| azure content safety | 78    | 20     | 58     | 25.6%  | 718    |
| azure prompt shields | 78    | 19     | 59     | 24.4%  | 1534   |
| aws bedrock          | 78    | 46     | 32     | 59.0%  | 807    |
| llama firewall       | 78    | 67     | 11     | 85.9%  | 51     |
| llm guard            | 78    | 67     | 11     | 85.9%  | 207    |

## Category Winners

| OWASP Ref | Category                  | Winner Backend  |
|-----------|---------------------------|-----------------|
| LLM01     | Prompt Injection          | llama firewall* |
| LLM02     | Insecure Output           | llama firewall* |
| LLM03     | Training Data Poisoning   | llama firewall* |
| LLM04     | Model DoS                 | llama firewall* |
| LLM05     | Supply Chain              | llama firewall* |
| LLM06     | Sensitive Info Disclosure | llama firewall* |
| LLM07     | Insecure Plugin           | llama firewall* |
| LLM08     | Excessive Agency          | llama firewall* |
| LLM09     | Overreliance              | llama firewall* |
| LLM10     | Model Theft               | llama firewall* |

- \* LLM01: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, lakera, openai moderation, azure prompt shields, aws bedrock, llama firewall, llm guard
- \* LLM02: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, lakera, llama firewall, llm guard
- \* LLM03: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, lakera, aws bedrock, llama firewall, llm guard
- \* LLM04: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, lakera, llama firewall, llm guard

- \* LLM05: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, llama firewall, llm guard
- \* LLM06: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, lakera, llama firewall, llm guard
- \* LLM07: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, llama firewall, llm guard
- \* LLM08: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, lakera, llama firewall, llm guard
- \* LLM09: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, lakera, llama firewall, llm guard
- \* LLM10: Tiebreaker by latency — llama firewall (51 ms) wins among: nemo, lakera, llama firewall, llm guard

## Independence Statement

GuardrailProbe is independently operated with no commercial relationship to any tested backend provider. NVIDIA, Microsoft, OpenAI, Lakera, Meta, Protect AI, and Amazon do not fund, endorse, or influence this project or its methodology.

Probe library, scoring logic, and methodology are fully open source and independently auditable at:  
**[github.com/askuma/guardrailprobe](https://github.com/askuma/guardrailprobe)**

This report was generated using probe library v0.1.0 against the backends listed above. Results reflect backend behavior at the time of testing with default or minimal configuration unless otherwise noted.

For methodology details see METHODOLOGY.md.

This document is digitally signed with an electronic signature and an RFC 3161 trusted timestamp. Verify the signature using any PDF viewer or via the GET `/redteam/reports/aaf34a3a-b850-4085-b01a-4fbcd70983ee/export` endpoint.